

Project2-Heart disease prediction

April 8, 2025

1 Project: Heart Disease Prediction Model using Logistic Regression

Objective of Project: To predict the presence of heart disease in patients using clinical data.

Dataset: The classic Heart Disease dataset (Cleveland dataset from UCI).

```
[3]: # import libraries
import pandas as pd
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings('ignore')
```

```
[5]: # load the dataset
df = pd.read_csv('heart.csv')
df.head()
```

```
[5]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
[9]: df.shape
```

```
[9]: (303, 14)
```

1.1 Data Dictionary (from domain expertise) age: age in years

sex: sex

1 = male

0 = female

cp: chest pain type

Value 0: typical angina

Value 1: atypical angina

Value 2: non-anginal pain

Value 3: asymptomatic

trestbps: resting blood pressure (in mm Hg on admission to the hospital)

chol: serum cholesterol in mg/dl

fbs: (fasting blood sugar > 120 mg/dl)

1 = true;

0 = false

restecg: resting electrocardiographic results

Value 0: normal

Value 1: having ST-T wave abnormality (T wave inversions and/or ST elevation or depression of > 0.05 mV)

Value 2: showing probable or definite left ventricular hypertrophy by Estes' criteria

thalach: maximum heart rate achieved

exang: exercise induced angina

1 = yes

0 = no

oldpeak = ST depression induced by exercise relative to rest

slope: the slope of the peak exercise ST segment

Value 0: upsloping

Value 1: flat

Value 2: downsloping

ca: number of major vessels (0-3) colored by fluoroscopy

thal:

0 = error (in the original dataset 0 maps to NaN's)

1 = fixed defect

2 = normal

3 = reversible defect

target (the 1

Note on the target label:

Diagnosis of heart disease (angiographic disease status) Value 0: < 50 Value 1: > 50 (able):

0 = no disease,

1 = disease

```
[13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    int64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    int64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
[15]: df.describe()
```

```
[15]:
```

	age	sex	cp	trestbps	chol	fbs	\
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	

	restecg	thalach	exang	oldpeak	slope	ca	\
--	---------	---------	-------	---------	-------	----	---

count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	0.528053	149.646865	0.326733	1.039604	1.399340	0.729373
std	0.525860	22.905161	0.469794	1.161075	0.616226	1.022606
min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000
50%	1.000000	153.000000	0.000000	0.800000	1.000000	0.000000
75%	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000
max	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000

	thal	target
count	303.000000	303.000000
mean	2.313531	0.544554
std	0.612277	0.498835
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

1.2 Data pre-processing

```
[24]: # check the missing values
```

```
df.isnull().sum()
```

```
[24]: age          0
sex          0
cp          0
trestbps    0
chol        0
fbs         0
restecg     0
thalach     0
exang       0
oldpeak     0
slope       0
ca          0
thal        0
target      0
dtype: int64
```

```
[28]: # split the data in independent and dependent variables
```

```
x= df.drop('target', axis = 1) # independent feature
y = df['target']               # Dependent feature (target variable)
```

```
[44]: # train test split
```

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.2,
↳random_state=42)
```

```
[46]: x_train
```

```
[46]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
132	42	1	1	120	295	0	1	162	0	0.0	
202	58	1	0	150	270	0	0	111	1	0.8	
196	46	1	2	150	231	0	1	147	0	3.6	
75	55	0	1	135	250	0	0	161	0	1.4	
176	60	1	0	117	230	1	1	160	1	1.4	
..	
188	50	1	2	140	233	0	1	163	0	0.6	
71	51	1	2	94	227	0	1	154	1	0.0	
106	69	1	3	160	234	1	0	131	0	0.1	
270	46	1	0	120	249	0	0	144	0	0.8	
102	63	0	1	140	195	0	1	179	0	0.0	

	slope	ca	thal
132	2	0	2
202	2	0	3
196	1	0	2
75	1	0	2
176	2	2	3
..
188	1	1	3
71	2	1	3
106	1	1	2
270	2	0	3
102	2	2	2

```
[242 rows x 13 columns]
```

```
[48]: y_train
```

```
[48]:
```

132	1
202	0
196	0
75	1
176	0
..	
188	0
71	1
106	1
270	0
102	1

```
Name: target, Length: 242, dtype: int64
```

```
[50]: from sklearn.linear_model import LogisticRegression

[52]: logreg = LogisticRegression()

logreg.fit(x_train, y_train)

[52]: LogisticRegression()

[54]: # Make the predictions
y_pred = logreg.predict(x_test)

[56]: y_pred

[56]: array([0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0,
          0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1,
          1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0], dtype=int64)

[58]: y_pred_prob = logreg.predict_proba(x_test)[: , 1]

[60]: y_pred_prob

[60]: array([1.17702855e-01, 7.89568007e-01, 8.11578819e-01, 2.59305607e-02,
          9.30504069e-01, 9.05120923e-01, 6.06542584e-01, 1.14143705e-03,
          4.83333573e-03, 5.47484409e-01, 7.89531849e-01, 7.11814629e-02,
          9.25950579e-01, 2.19235026e-02, 9.82205572e-01, 9.52954802e-01,
          9.77178756e-01, 5.04351348e-02, 6.34987408e-03, 1.02764744e-02,
          7.38387361e-01, 1.05222577e-02, 1.37311576e-01, 8.00450586e-01,
          9.19348268e-01, 6.67266260e-01, 9.08248067e-01, 6.86177887e-01,
          6.56726108e-03, 9.10155769e-01, 4.21340094e-02, 2.93431981e-02,
          5.62116979e-03, 7.24593241e-02, 6.90302723e-01, 6.84994624e-02,
          6.38464062e-01, 8.73556106e-01, 8.09302006e-01, 8.53214445e-01,
          5.17951451e-01, 8.35213292e-01, 8.15581972e-01, 6.86176207e-01,
          8.41197991e-01, 5.46301117e-03, 8.01275776e-01, 9.52859388e-01,
          7.11805039e-02, 3.04629283e-02, 6.34226202e-02, 1.25017980e-02,
          8.70138566e-01, 9.77604269e-01, 2.22735234e-01, 5.80512878e-04,
          4.32857695e-02, 9.63002856e-01, 1.00577902e-02, 3.81726863e-03,
          2.93434520e-02])

[62]: # Evalauation of the logistic regression

from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

[64]: print('Accuracy score is:', accuracy_score(y_test, y_pred))

Accuracy score is: 0.8852459016393442
```

```
[66]: print('Confusion Matrix')
      print(confusion_matrix(y_test, y_pred))
```

```
Confusion Matrix
[[25  4]
 [ 3 29]]
```

```
[68]: print('classification Report')
      print(classification_report(y_test, y_pred))
```

```
classification Report
```

	precision	recall	f1-score	support
0	0.89	0.86	0.88	29
1	0.88	0.91	0.89	32
accuracy			0.89	61
macro avg	0.89	0.88	0.88	61
weighted avg	0.89	0.89	0.89	61

```
[72]: import joblib

      # Save your trained model ( Logistic Regression)
      joblib.dump(logreg, "heart_disease_model.pkl")
```

```
[72]: ['heart_disease_model.pkl']
```

This project successfully demonstrates the application of machine learning techniques to predict the likelihood of heart disease based on patient attributes. Among the models tested, [logreg] performed the best with an accuracy of 88%. This indicates that the model can assist healthcare professionals in early screening of high-risk individuals, potentially improving patient outcomes through timely interventions.

2 Thank You!!

```
[ ]:
```